

Uncertainty in Machine Learning \ Assignment 2

Massimiliano Falzari (s3459101)

June 9, 2022

Contents

| | | |
|----------|--|----------|
| 1 | Is the ELBO a proper scoring rule? | 1 |
| 2 | Disentanglement Applications | 2 |
| 3 | Programming Disentangling Uncertainty | 2 |

1 Is the ELBO a proper scoring rule?

Yes, the ELBO is a proper scoring rule. By definition a proper scoring rule has the following property:

$$S(p_\theta, q) \leq S(q, q)$$

With equality only holding with $p_\theta(y|x) = q(y|x)$.

We know that the KL-divergence can be used as a scoring rule if we define it as follows:

$$S(p_\theta, q) = -D_{KL}(p_\theta, q)$$

we know that the KL-divergence is always non-negative and $D_{KL}(p, q) = 0$ iff $p = q$ Therefore it satisfies the aforementioned property

Finally, we can rewrite the problem of minimizing the KL-divergence as maximizing the ELBO. (derivation in the slides)

Therefore, the ELBO is equivalent to the - KL-divergence up to a constant which means that it is a proper scoring rule.

2 Disentanglement Applications

- Spam detection

In this case, in order to decide whether it is spam or not, we only want to focus on epistemic uncertainty, and we want to remove the aleatoric uncertainty. It is a case of anomaly detection.

- Forecasting/Planning

When performing forecasting/planning, it can be quite important to distinguish between epistemic and aleatoric uncertainty. In particular, it can be useful because we can decide how to weight these two types of uncertainty. For example, let's say we have two or more predictions/forecasts/plans with the same total uncertainty. In this case, it can be crucial to distinguish between aleatoric and epistemic uncertainty. We probably want to favour plans that have high aleatoric uncertainty as opposed to high epistemic uncertainty (Note we are at inference time).

3 Programming Disentangling Uncertainty

I have used an Ensemble as an uncertainty quantification method. In particular, a 5-network ensemble was used. As we can see the domain of the training data is roughly between 0.5 and -0.5 mean while the testing domain is roughly between 1 and -1. This was done in order to verify whether the epistemic uncertainty will increase or not outside of the training domain. As we can see from Figure 5, after we go outside of the training domain, the epistemic uncertainty is higher than when we are inside the training domain. Therefore, this shows that the estimate seems reasonable (Even though, I expected a higher value).

For the aleatoric uncertainty, on the other hand, (Figure 4) the network failed to fit the function we used as noise. It almost looks like the noise is fixed and is not based on the input. That said, even though the shape of the noise function was not captured, the value seems to be in line with the mean of the function which is roughly 0.5.

Lastly, we can see that the predicted values as a similar behaviour of the epistemic uncertainty. In particular, the more we go outside of the training domain, the more the predicted values does not follow the function we wanted to fit.(Figure 3)

Finally, the implementation used the DeepEnsembleRegressor from the `keras_uncertainty` library. Generally, an ensemble is simply a set of networks

that are trained on the same data. At inference time (this was the main feature used from the library) the ensemble estimates from each of the networks the mean and variance predicted and it returns the mean of the means and the mean of the variance. In our case, we used the flag `disentangle_uncertainty=True` which instead of returning the mean and variance of the ensemble returns the standard deviation of the means for the epistemic uncertainty and the mean of the standard deviation for the aleatoric uncertainty. These were the main features used from the library.




Figure 1: training data




Figure 2: testing data




Figure 3: predicted

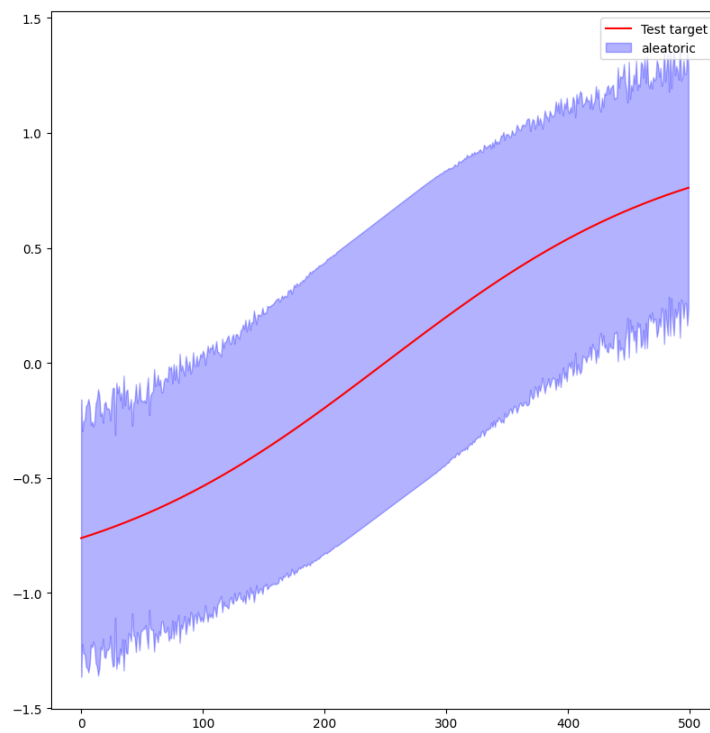


Figure 4: aleatoric uncertainty




Figure 5: epistemic uncertainty